



**Vidyavardhini's College of Engineering and Technology**

**Department of Artificial Intelligence & Data Science**

Experiment No. 5
Implementation of Expectation-Maximization Algorithm
Date of Performance:
Date of Submission:
Marks:
Sign:



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

**Aim:** Implementation of Expectation-Maximization Algorithm.

**Objective:** Able to interpret the how Expectation Maximization Algorithm is applied in unsupervised clustering algorithm to handle the non observable features.

### Theory:


In the real-world applications of machine learning, it is very common that there are many relevant features available for learning but only a small subset of them are observable. So, for the variables which are sometimes observable and sometimes not, then we can use the instances when that variable is visible is observed for the purpose of learning and then predict its value in the instances when it is not observable.

On the other hand, *Expectation-Maximization algorithm* can be used for the latent variables (variables that are not directly observable and are actually inferred from the values of the other observed variables) too in order to predict their values with the condition that the general form of probability distribution governing those latent variables is known to us. This algorithm is actually at the base of many unsupervised clustering algorithms in the field of machine learning.

It was explained, proposed and given its name in a paper published in 1977 by Arthur Dempster, Nan Laird, and Donald Rubin. It is used to find the *local maximum likelihood parameters* of a statistical model in the cases where latent variables are involved and the data is missing or incomplete.

### Algorithm:

1. Given a set of incomplete data, consider a set of starting parameters.
2. **Expectation step (E – step):** Using the observed available data of the dataset, estimate (guess) the values of the missing data.
3. **Maximization step (M – step):** Complete data generated after the expectation (E) step is used in order to update the parameters.
4. Repeat step 2 and step 3 until convergence.

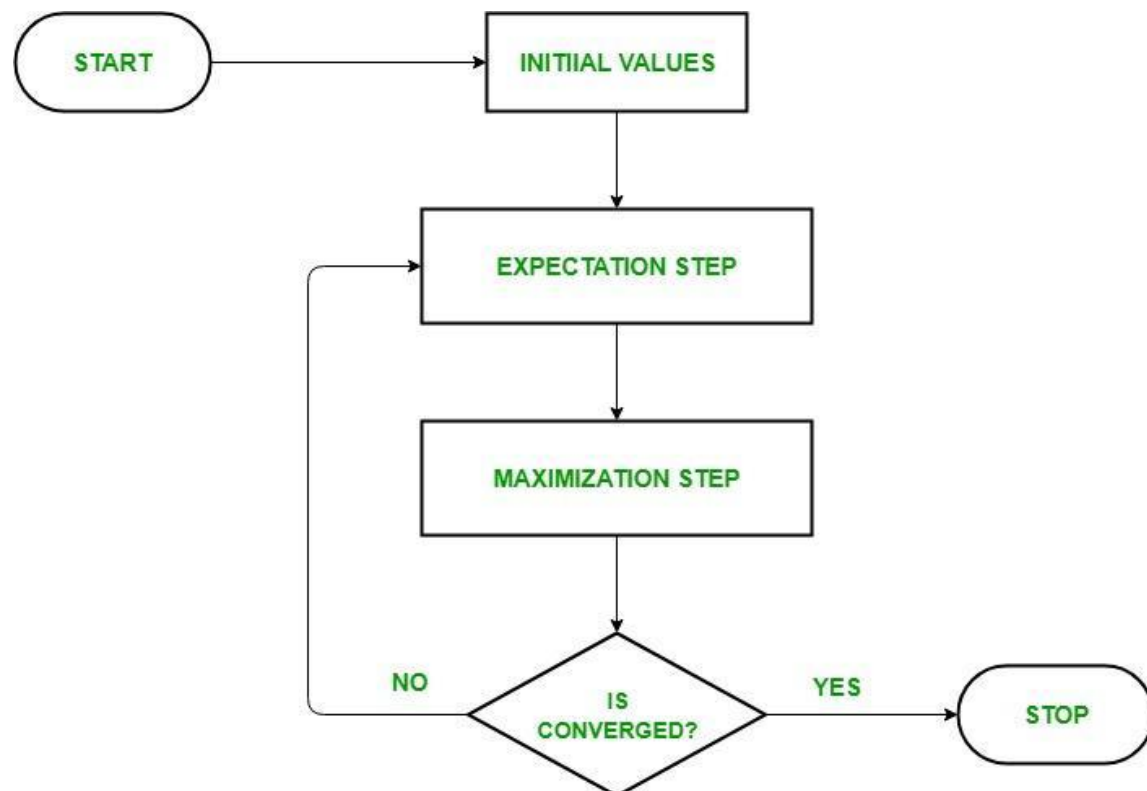




The essence of Expectation-Maximization algorithm is to use the available observed data of the dataset to estimate the missing data and then using that data to update the values of the parameters.

- Initially, a set of initial values of the parameters are considered. A set of incomplete observed data is given to the system with the assumption that the observed data comes from a specific model.
- The next step is known as “Expectation” – step or *E-step*. In this step, we use the observed data in order to estimate or guess the values of the missing or incomplete data. It is basically used to update the variables.
- The next step is known as “Maximization”-step or *M-step*. In this step, we use the complete data generated in the preceding “Expectation” – step in order to update the values of the parameters. It is basically used to update the hypothesis.
- Now, in the fourth step, it is checked whether the values are converging or not, if yes, then stop otherwise repeat *step-2* and *step-3* i.e. “Expectation” – step and “Maximization” – step until the convergence occurs.

**Flow chart for EM algorithm –**





# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

### Usage of EM algorithm –

- It can be used to fill the missing data in a sample.
- It can be used as the basis of unsupervised learning of clusters.
- It can be used for the purpose of estimating the parameters of Hidden Markov Model (HMM).
- It can be used for discovering the values of latent variables.

### Advantages of EM algorithm –

- It is always guaranteed that likelihood will increase with each iteration.
- The E-step and M-step are often pretty easy for many problems in terms of implementation.
- Solutions to the M-steps often exist in the closed form.

### Disadvantages of EM algorithm –

- It has slow convergence.
- It makes convergence to the local optima only.
- It requires both the probabilities, forward and backward (numerical optimization requires only forward probability).

### Implementation:

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style("white")
%matplotlib inline
#for matrix math
import numpy as np
#for normalization + probability density function computation
from scipy import stats
#for data preprocessing
import pandas as pd
from math import sqrt, log, exp, pi
from random import uniform
print("import done")
```

```
import done
```



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
random_seed=36788765
```

```
np.random.seed(random_seed)
```

```
Mean1 = 2.0 # Input parameter, mean of first normal probability distribution
```

```
Standard_dev1 = 4.0 #@param {type:"number"}
```

```
Mean2 = 9.0 # Input parameter, mean of second normal probability distribution
```

```
Standard_dev2 = 2.0 #@param {type:"number"}
```

```
# generate data
```

```
y1 = np.random.normal(Mean1, Standard_dev1, 1000)
```

```
y2 = np.random.normal(Mean2, Standard_dev2, 500)
```

```
data=np.append(y1,y2)
```

```
# For data visualisation calculate left and right of the graph
```

```
Min_graph = min(data)
```

```
Max_graph = max(data)
```

```
x = np.linspace(Min_graph, Max_graph, 2000) # to plot the data
```

```
print('Input Gaussian {:}:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ '.format("1", Mean1, Standard_dev1))
```

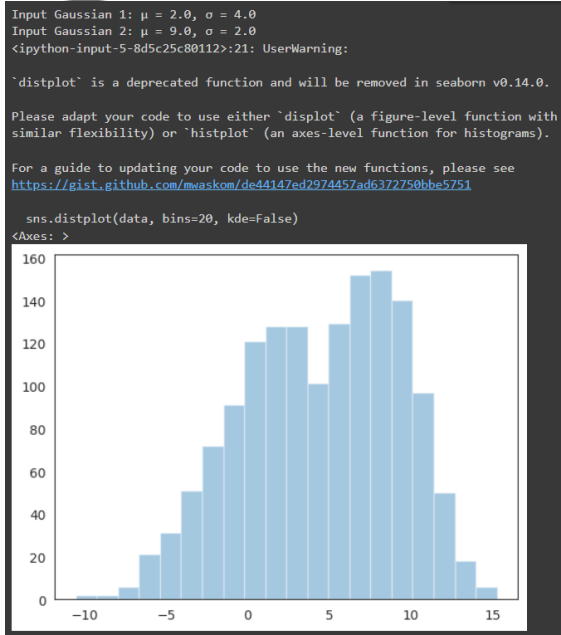
```
print('Input Gaussian {:}:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ '.format("2", Mean2, Standard_dev2))
```

```
sns.distplot(data, bins=20, kde=False)
```



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science



class Gaussian:

"Model univariate Gaussian"

def \_\_init\_\_(self, mu, sigma):

#mean and standard deviation

self.mu = mu

self.sigma = sigma

#probability density function

def pdf(self, datum):

"Probability of a data point given the current parameters"

u = (datum - self.mu) / abs(self.sigma)

y = (1 / (sqrt(2 \* pi) \* abs(self.sigma))) \* exp(-u \* u / 2)

return y

def \_\_repr\_\_(self):

return 'Gaussian({0:4.6}, {1:4.6})'.format(self.mu, self.sigma)

print("done")

done




# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
best_single = Gaussian(np.mean(data), np.std(data))
print('Best single Gaussian:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ '.format(best_single.mu, best_single.sigma))
#fit a single gaussian curve to the data
g_single = stats.norm(best_single.mu, best_single.sigma).pdf(x)
sns.distplot(data, bins=20, kde=False, norm_hist=True)
plt.plot(x, g_single, label='single gaussian')
plt.legend()
```

```
Best single Gaussian:  $\mu = 4.4$ ,  $\sigma = 4.8$ 
<ipython-input-7-0e8931e2d352>:5: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).
For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372759bbe5751
sns.distplot(data, bins=20, kde=False, norm_hist=True)
<matplotlib.legend.Legend at 0x7f718abe36d0>
```



```
class GaussianMixture_self:
```

```
    "Model mixture of two univariate Gaussians and their EM estimation"
```

```
    def __init__(self, data, mu_min=min(data), mu_max=max(data), sigma_min=1,
sigma_max=1, mix=.5):
```

```
        self.data = data
```

```
        #todo the Algorithm would be numerical enhanced by normalizing the data first, next do
all the EM steps and do the de-normalising at the end
```

```
        #init with multiple gaussians
```

```
        self.one = Gaussian(uniform(mu_min, mu_max),
```



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
uniform(sigma_min, sigma_max))

self.two = Gaussian(uniform(mu_min, mu_max),
                    uniform(sigma_min, sigma_max))

#as well as how much to mix them
self.mix = mix

def Estep(self):
    "Perform an E(stimation)-step, assign each point to gaussian 1 or 2 with a percentage"
    # compute weights
    self.loglike = 0. # = log(p = 1)
    for datum in self.data:
        # unnormalized weights
        wp1 = self.one.pdf(datum) * self.mix
        wp2 = self.two.pdf(datum) * (1. - self.mix)
        # compute denominator
        den = wp1 + wp2
        # normalize
        wp1 /= den
        wp2 /= den    # wp1+wp2= 1, it either belongs to gaussian 1 or gaussian 2
        # add into loglike
        self.loglike += log(den) #freshening up self.loglike in the process
        # yield weight tuple
        yield (wp1, wp2)

def Mstep(self, weights):
    "Perform an M(aximization)-step"
    # compute denominators
    (left, right) = zip(*weights)
    one_den = sum(left)
    two_den = sum(right)
```





# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
# compute new means
self.one.mu = sum(w * d for (w, d) in zip(left, data)) / one_den
self.two.mu = sum(w * d for (w, d) in zip(right, data)) / two_den

# compute new sigmas
self.one.sigma = sqrt(sum(w * ((d - self.one.mu) ** 2)
                        for (w, d) in zip(left, data)) / one_den)
self.two.sigma = sqrt(sum(w * ((d - self.two.mu) ** 2)
                        for (w, d) in zip(right, data)) / two_den)

# compute new mix
self.mix = one_den / len(data)

def iterate(self, N=1, verbose=False):
    "Perform N iterations, then compute log-likelihood"
    for i in range(1, N+1):
        self.Mstep(self.Estep()) #The heart of the algorithm, perform E-step and next M-step
        if verbose:
            print('{0:2} {1}'.format(i, self))
        self.Estep() # to freshen up self.loglike

def pdf(self, x):
    return (self.mix)*self.one.pdf(x) + (1-self.mix)*self.two.pdf(x)

def __repr__(self):
    return 'GaussianMixture({0}, {1}, mix={2.03})'.format(self.one,
                                                         self.two,
                                                         self.mix)

def __str__(self):
    return 'Mixture: {0}, {1}, mix={2:.03}'.format(self.one,
                                                  self.two,
                                                  self.mix)
```



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
print("done")
```

done

```
n_iterations = 20
```

```
best_mix = None
```

```
best_loglike = float('-inf')
```

```
mix = GaussianMixture_self(data)
```

```
for _ in range(n_iterations):
```

```
    try:
```

```
        #train!
```

```
        mix.iterate(verbose=True)
```

```
        if mix.loglike > best_loglike:
```

```
            best_loglike = mix.loglike
```

```
            best_mix = mix
```

```
    except (ZeroDivisionError, ValueError, RuntimeWarning): # Catch division errors from
```

```
bad starts, and just throw them out...
```

```
        print("one less")
```

```
    pass
```

```
1 Mixture: Gaussian(-5.81052, 1.27375), Gaussian(4.83277, 4.39033), mix=0.036)
1 Mixture: Gaussian(-5.30703, 1.23684), Gaussian(4.73186, 4.51064), mix=0.0281)
1 Mixture: Gaussian(-5.04445, 1.20674), Gaussian(4.7003, 4.54811), mix=0.0257)
1 Mixture: Gaussian(-4.88152, 1.20146), Gaussian(4.68792, 4.56395), mix=0.0249)
1 Mixture: Gaussian(-4.76296, 1.21164), Gaussian(4.68347, 4.57134), mix=0.0247)
1 Mixture: Gaussian(-4.66629, 1.23092), Gaussian(4.6833, 4.57443), mix=0.025)
1 Mixture: Gaussian(-4.58114, 1.25563), Gaussian(4.68588, 4.57496), mix=0.0255)
1 Mixture: Gaussian(-4.50225, 1.28362), Gaussian(4.69047, 4.57376), mix=0.0262)
1 Mixture: Gaussian(-4.42685, 1.31357), Gaussian(4.69664, 4.5713), mix=0.0271)
1 Mixture: Gaussian(-4.35343, 1.3447), Gaussian(4.70413, 4.56787), mix=0.0281)
1 Mixture: Gaussian(-4.28121, 1.3765), Gaussian(4.71277, 4.56364), mix=0.0292)
1 Mixture: Gaussian(-4.20982, 1.40866), Gaussian(4.72243, 4.55874), mix=0.0305)
1 Mixture: Gaussian(-4.13909, 1.44099), Gaussian(4.73304, 4.55326), mix=0.0319)
1 Mixture: Gaussian(-4.06903, 1.47337), Gaussian(4.74451, 4.54725), mix=0.0334)
1 Mixture: Gaussian(-3.99968, 1.50573), Gaussian(4.75682, 4.54077), mix=0.0351)
1 Mixture: Gaussian(-3.93113, 1.53806), Gaussian(4.7699, 4.53384), mix=0.0368)
1 Mixture: Gaussian(-3.86347, 1.57032), Gaussian(4.78374, 4.5265), mix=0.0386)
1 Mixture: Gaussian(-3.79679, 1.60253), Gaussian(4.79829, 4.51876), mix=0.0405)
1 Mixture: Gaussian(-3.73116, 1.63468), Gaussian(4.81353, 4.51064), mix=0.0426)
1 Mixture: Gaussian(-3.66663, 1.6668), Gaussian(4.82946, 4.50216), mix=0.0447)
```



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
n_iterations = 300
n_random_restarts = 4
best_mix = None
best_loglike = float('-inf')
print('Computing best model with random restarts...\n')
for _ in range(n_random_restarts):
    mix = GaussianMixture_self(data)
    for _ in range(n_iterations):
        try:
            mix.iterate()
            if mix.loglike > best_loglike:
                best_loglike = mix.loglike
                best_mix = mix
        except (ZeroDivisionError, ValueError, RuntimeWarning): # Catch division errors from
            # bad starts, and just throw them out...
            pass
    #print('Best Gaussian Mixture :  $\mu = \{:.2\}$ ,  $\sigma = \{:.2\}$  with  $\mu = \{:.2\}$ ,  $\sigma =$ 
    # $\{:.2\}$ '.format(best_mix.one.mu, best_mix.one.sigma, best_mix.two.mu, best_mix.two.sigma))

print('Input Gaussian {::}:  $\mu = \{:.2\}$ ,  $\sigma = \{:.2\}$ '.format("1", Mean1, Standard_dev1))
print('Input Gaussian {::}:  $\mu = \{:.2\}$ ,  $\sigma = \{:.2\}$ '.format("2", Mean2, Standard_dev2))
print('Gaussian {::}:  $\mu = \{:.2\}$ ,  $\sigma = \{:.2\}$ , weight =  $\{:.2\}$ '.format("1", best_mix.one.mu,
best_mix.one.sigma, best_mix.mix))
print('Gaussian {::}:  $\mu = \{:.2\}$ ,  $\sigma = \{:.2\}$ , weight =  $\{:.2\}$ '.format("2", best_mix.two.mu,
best_mix.two.sigma, (1-best_mix.mix)))
#Show mixture
sns.distplot(data, bins=20, kde=False, norm_hist=True)
g_both = [best_mix.pdf(e) for e in x]
plt.plot(x, g_both, label='gaussian mixture')
g_left = [best_mix.one.pdf(e) * best_mix.mix for e in x]
plt.plot(x, g_left, label='gaussian one')
g_right = [best_mix.two.pdf(e) * (1-best_mix.mix) for e in x]
```

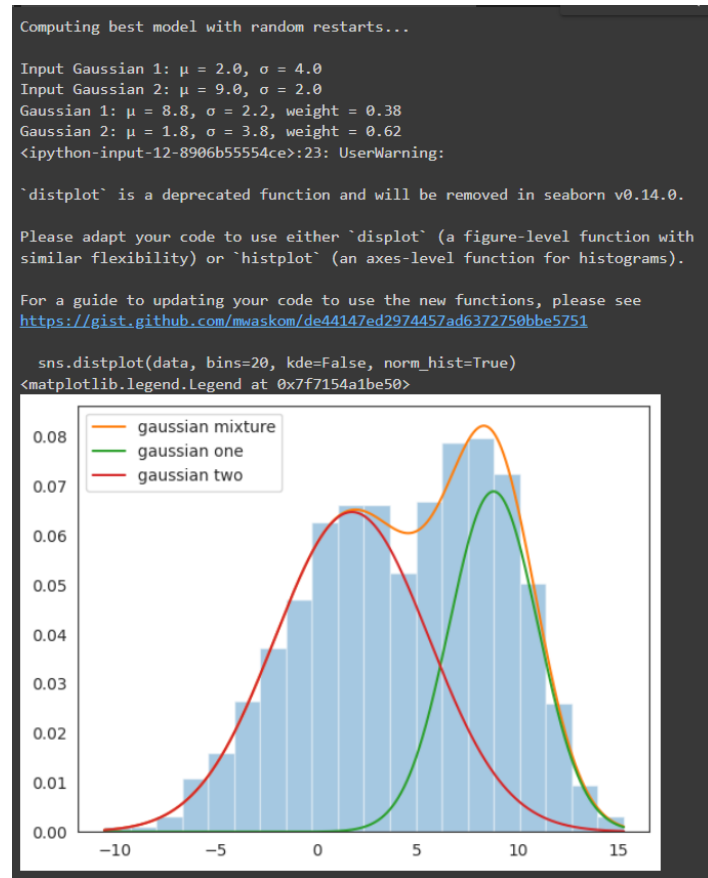


# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

```
plt.plot(x, g_right, label='gaussian two')
```

```
plt.legend()
```



```
from sklearn.mixture import GaussianMixture
```

```
gmm = GaussianMixture(n_components = 2, tol=0.000001)
```

```
gmm.fit(np.expand_dims(data, 1)) # Parameters: array-like, shape (n_samples, n_features), 1
dimension dataset so 1 feature
```

```
Gaussian_nr = 1
```

```
print('Input Gaussian {}:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ '.format("1", Mean1, Standard_dev1))
```

```
print('Input Gaussian {}:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ '.format("2", Mean2, Standard_dev2))
```

```
for mu, sd, p in zip(gmm.means_.flatten(), np.sqrt(gmm.covariances_.flatten()),
```

```
gmm.weights_):
```

```
    print('Gaussian {}:  $\mu = {:.2}$ ,  $\sigma = {:.2}$ , weight = {:.2}'.format(Gaussian_nr, mu, sd, p))
```

```
    g_s = stats.norm(mu, sd).pdf(x) * p
```

```
    plt.plot(x, g_s, label='gaussian sklearn')
```

```
    Gaussian_nr += 1
```



# Vidyavardhini's College of Engineering and Technology

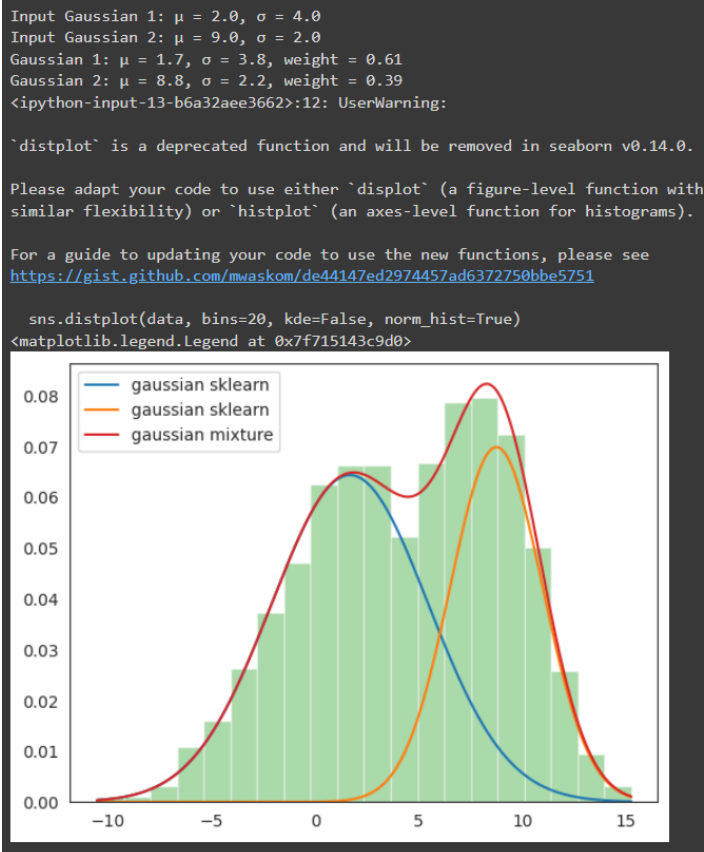
## Department of Artificial Intelligence & Data Science

```
sns.distplot(data, bins=20, kde=False, norm_hist=True)
```

```
gmm_sum = np.exp([gmm.score_samples(e.reshape(-1, 1)) for e in x]) #gmm gives log probability, hence the exp() function
```

```
plt.plot(x, gmm_sum, label='gaussian mixture')
```

```
plt.legend()
```



### Conclusion:

- 1) What is the use of Expectation-Maximization Algorithm?

The Expectation-Maximization (EM) algorithm is a powerful statistical method used for estimating parameters in probabilistic models when dealing with incomplete or missing data. It iteratively alternates between two main steps: the E-step (Expectation), where it calculates the expected values of latent variables given the observed data and current parameter estimates, and the M-step (Maximization), where it updates the parameters based on the expected values obtained in the E-step. EM is particularly useful in scenarios where data is partially observed or latent variables are involved, such as in clustering algorithms like Gaussian Mixture Models (GMM) or in fitting models with hidden variables like Hidden Markov Models (HMM). Its



# Vidyavardhini's College of Engineering and Technology

## Department of Artificial Intelligence & Data Science

versatility and ability to handle missing data make it invaluable in various fields, including machine learning, bioinformatics, and natural language processing.

### 2) How it can be used to handle missing data?

The Expectation-Maximization (EM) algorithm is particularly useful for handling missing data because of its ability to iteratively estimate parameters in probabilistic models while accommodating incomplete or partially observed data. Here's how it can be used to handle missing data:

1)E-step (Expectation): In the E-step of the EM algorithm, the missing data points are treated as latent variables. The algorithm calculates the expected values of these missing data points based on the observed data and the current estimates of the model parameters.

2)M-step (Maximization): In the M-step, the algorithm updates the parameters of the model based on the expected values obtained in the E-step, along with the observed data. This step involves maximizing the likelihood function, taking into account both the observed and imputed values.

3)Iterative Process: The EM algorithm iteratively alternates between the E-step and M-step until convergence, where the parameters stabilize. During each iteration, the algorithm refines its estimates of both the model parameters and the missing data values.

4)Imputation of Missing Values: Through this iterative process, the EM algorithm effectively imputes missing values by estimating their most likely values based on the observed data and the model parameters. These imputed values can then be used for subsequent analyses or modeling tasks.